

Parallel Programming for FPGAs (part 1) Practice

Lab1: Finite impulse response (FIR) filter

- 有限冲激响应滤波器

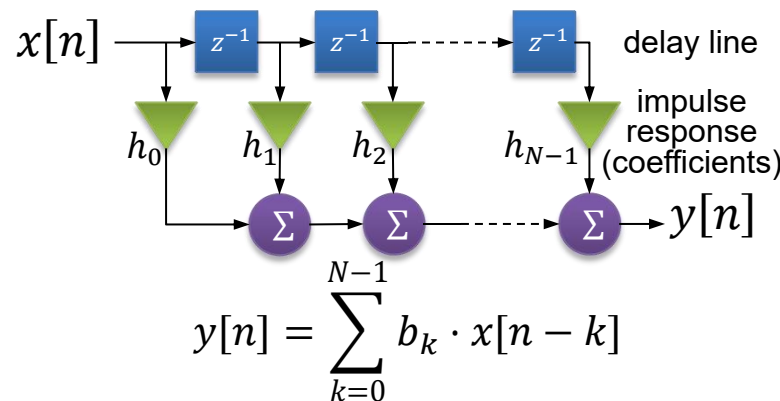
- $$y[n] = b_0x[n] + b_1x[n - 1] + \cdots + b_Nx[n - N]$$

$$= \sum_{i=0}^N b_i \cdot x[n - i] \text{ (discrete convolution)}$$

- $x[n]$ is the input signal, $y[n]$ is the output signal

- N is the filter order

- b_i is the impulse response at the i th instant (*coefficient*)



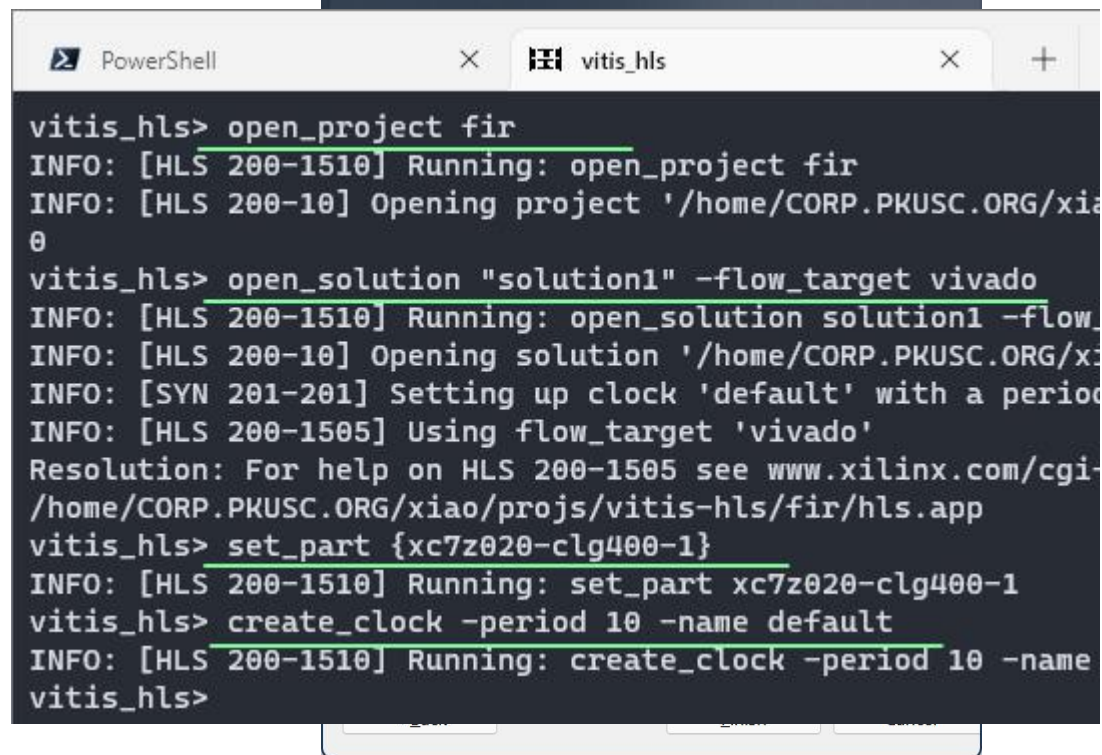
Create FIR project (1)

GUI

- Open *Vitis HLS* GUI
- Click `Create Project`
- Set project name
- Click `Next`s
 - until `Solution config.`
- Set Clock Period: 10 (ns)
- Set Part: xc7z020clg400-1

Cmd

- `source dir-to-vitis-hls/settings64.sh`
- `vitis_hls -i`: enter interactive mode
- type commands



```
PowerShell vitis_hls
vitis_hls> open_project fir
INFO: [HLS 200-1510] Running: open_project fir
INFO: [HLS 200-10] Opening project '/home/CORP.PKUSC.ORG/xiao
0
vitis_hls> open_solution "solution1" -flow_target vivado
INFO: [HLS 200-1510] Running: open_solution solution1 -flow.
INFO: [HLS 200-10] Opening solution '/home/CORP.PKUSC.ORG/xs
INFO: [SYN 201-201] Setting up clock 'default' with a period
INFO: [HLS 200-1505] Using flow_target 'vivado'
Resolution: For help on HLS 200-1505 see www.xilinx.com/cgi-
/home/CORP.PKUSC.ORG/xiao/projs/vitis-hls/fir/hls.app
vitis_hls> set_part {xc7z020-clg400-1}
INFO: [HLS 200-1510] Running: set_part xc7z020-clg400-1
vitis_hls> create_clock -period 10 -name default
INFO: [HLS 200-1510] Running: create_clock -period 10 -name
vitis_hls>
```

Create FIR project (2)


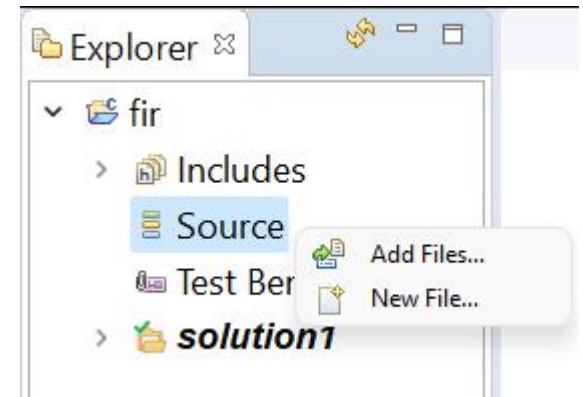
GUI

- Add Source files
 - fir.h, fir.cpp
- Add Test Bench
 - *.cpp, *.dat

CMD

- type commands

```
add_files fir_srcs/fir.cpp Source
add_files fir_srcs/fir.h
add_files -tb fir_srcs/fir.cpp
add_files -tb fir_srcs/fir_tb.cpp
add_files -tb fir_srcs/input.dat
add_files -tb fir_srcs/out.gold.dat
           test bench
```



Program design

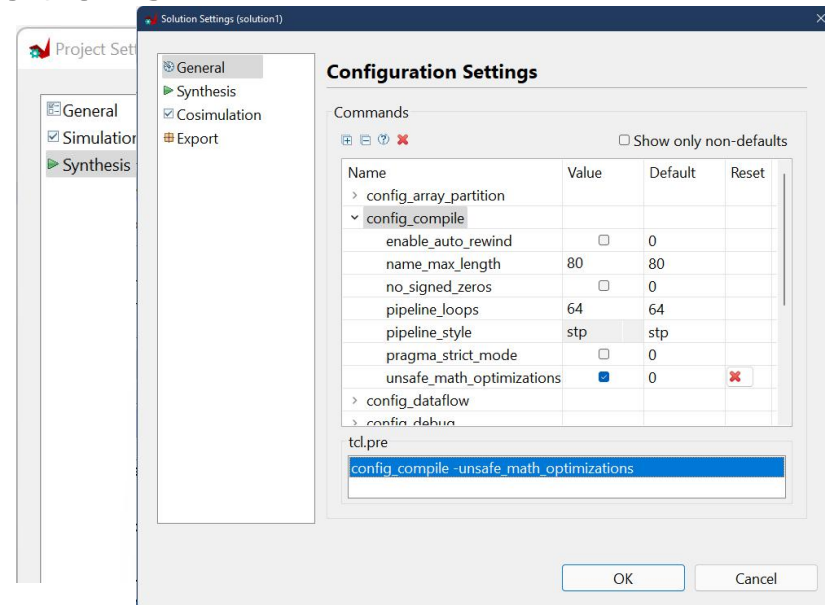

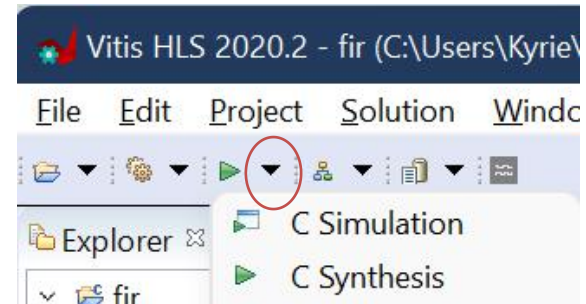
- *fir.h*
 - macro, typedef
 - function declaration
- *fir.cpp*
 - function definition
- *fir_tb.cpp*
 - ground truth
 - kernel call

```
// fir.h
#define N 11
typedef int coef_t;
#include "fir.h"
typedef int data_t;
void fir(data_t *y, data_t x) {
    typedef int acc_t;
    coef_t c[N] = {53, 0, -91, 0, 313,
500, 313, 0, -91, 0, 53};
    void fir(data_t *y, data_t x);
    static data_t shift_reg[N];
    acc_t acc = 0;
Shift_Accum_Loop:
    for (int i=N-1; i>=0; i--) {
        if (i==0) {
            acc += x * c[0];
            shift_reg[0] = x;
        } else {
            shift_reg[i] = shift_reg[i-1];
            acc += shift_reg[i] * c[i];
        }
    }
    *y = acc;
}
```

C-Simulation, C-Synthesis (1)

GUI

- Run C-Simulation
- Set top function
- Set synthesis configuration
 - enable unsafe math optimizations



- Run C Synthesis

C-Simulation, C-Synthesis (2)

CMD

- type commands

`csim_design`

`set_top fir`

`config_compile -unsafe_math_optimizations`

`csynth_design`

Observe reports (1)

GUI

- Open *Analysis* view – right top
- Look at *Synthesis Summary Report*

Debug Synthesis Analysis

GUI Screenshot: Synthesis Summary Report of 'fir'

General Information:

- Date: Thu Mar 3 10:53:04 2022
- Version: 2020.2 (Build 3064766 on Wed Nov 18 09:12:45 MST 2020)
- Project: fir
- Solution: solution1 (Vivado IP Flow Target)
- Product family: zynq
- Target device: xc7z020-clg400-1

Timing Estimate:

Target	Estimated	Uncertainty
10.00 ns	6.912 ns	2.70 ns

Performance & Resource Estimates

Performance and Resource

latency

interval

Modules & Loops	Issue Type	Slack	Latency(cycles)	Latency(ns)	Iteration Latency	Interval	Trip Count	Pipelined	BRAM	DSP	FF	LUT	URAM
fir	II Violation	-	28	280.000	-	29	-	no	0	1	709	316	0
Shift_Accum_Loop	II Violation	-	25	250.000	5	2	11	yes	-	-	-	-	-

Observe reports (2)

CMD

- cd to the vitis-hls project directory
- cd to the solution directory
- cd to `syn` directory
- cd to `report` directory
- open `[top function]_csynth.rpt`

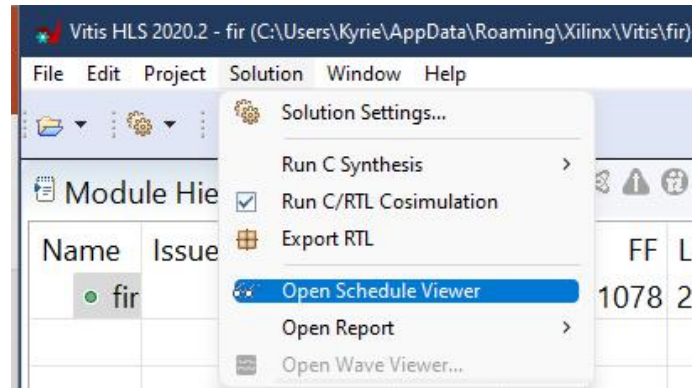
```
=====
== Vitis HLS Report for 'fir'
=====
* Date:                Sun Mar 13 15:52:19 2022

* Version:             2020.2 (Build 3064766 on Wed Nov 18 09:12:47 MST 2020)
* Project:             fir_hls_proj
* Solution:            solution1 (Vivado IP Flow Target)
* Product family:      zynq
* Target device:       xc7z020-clg400-1
```

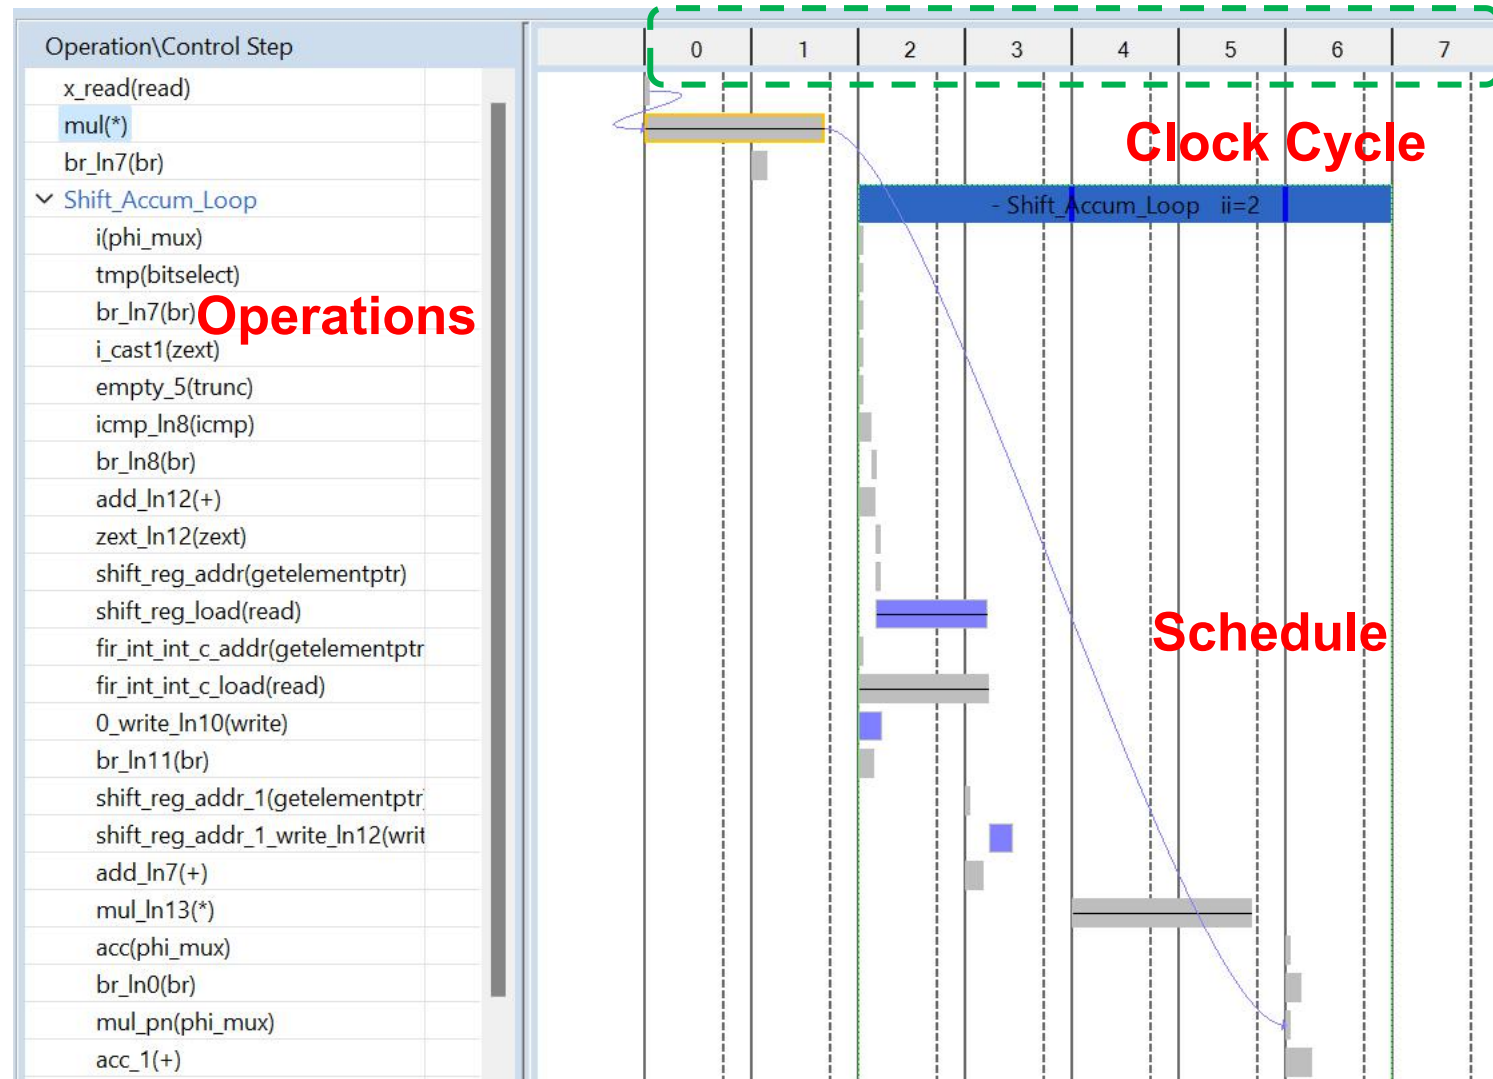
Observe Schedule Viewer (1)

GUI

- Open Schedule Viewr



Observe Schedule Viewer (2)



Observe Schedule Viewer (3)

CMD

- cd to solution directory
- cd to ``.autopilot/db/`` directory
- open ``[top_function].verbose.sched.rpt``
 - FSM

State

```
State 4! <SV = 3> <Delay = 6.91>
ST_4 : Operation 56 [2/2] (6.91ns) ----> "%mul_ln17 = mul
i32 %fir_int_int_c_load, i32 %UnifiedRetVal_i" [fir_srcs/fir.
cpp:17] ----> Operation 56 'mul' 'mul_ln17' <Predicate = (
!icmp_ln16)> <Delay = 6.91> <CoreInst = "Multiplier"> ---->
Core 3 'Multiplier' <Latency = 1> <II = 1> <Delay = 6.91> <
H/W Unit <Opcode : 'mul'> <InPorts = 2> <OutPorts = 1>
```

Operation

Lab2: Matrix Vector Multiplication

- MVM: $V_{out}[i] = \sum_k M[i][k] * V_{in}[k]$

$$\begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \odot \begin{bmatrix} x \\ y \\ z \end{bmatrix} = \begin{bmatrix} 1 \cdot x + 0 \cdot x + 0 \cdot x \\ 0 \cdot y + 1 \cdot y + 0 \cdot y \\ 0 \cdot z + 0 \cdot z + 1 \cdot z \end{bmatrix} = \begin{bmatrix} x \\ y \\ z \end{bmatrix}$$

Create baseline

- Please create and configure the project as you did in FIR example.
- Implement a simple baseline.
- Pass C-Simulation and run Synthesis

```
// mvm.h
typedef float in_t;
typedef float out_t;
#define N 64

// mvm.cpp
void mvm(in_t M[N][N], in_t V_In[N], out_t V_Out[N]) {
out_loop:
    for (int i = 0; i < N; i++) {
        out_t sum = 0;
dot_product_loop:
        for (int j = 0; j < N; j++) {
            sum += V_In[j] * M[i][j];
        }
        V_Out[i] = sum;
    }
}
```

Analyze MVM Baseline (1)

- What optimizations does *Vitis HLS* apply automatically?
 - Look at the *Vitis HLS Console* or *log information*.

```
INFO: [XFORM 203-510] Pipelining loop 'out_loop' (mvm_srcs/mvm.cpp:28) in function 'mvm' automatically.  
INFO: [XFORM 203-502] Unrolling all sub-loops inside loop 'out_loop' (mvm_srcs/mvm.cpp:28) in function 'mvm' for pipelining.  
INFO: [HLS 200-489] Unrolling loop 'dot_product_loop' (mvm_srcs/mvm.cpp:28) in function 'mvm' completely with a factor of 64.
```

- Pipeline “out_loop”, fully unroll “dot_product_loop”


Analyze MVM Baseline (2)

- Why the II is 32?

Name	Pipelined	Latency	Iteration Latency	Initiation Interval	Trip count
▼ • mvm	-	2376	-	2377	-
• out_loop	yes	2342	327	32	64

– Look at the *Schedule Viewer*

Analyze MVM Baseline (3)




- It take 32 cycles to read `V_In[]` and `M[i][]`

Analyze MVM Baseline (4)

- Why two *fadds* and two *fmuls*?

Name	BRAM	DSP	FF	LUT
▼ mvm	0	10	7547	4157
> I/O Ports(3)				
▼ Instances(4)	0	10	696	1422
♦ fadd_32ns_3i	0	2	205	390
♦ fadd_32ns_3i	0	2	205	390
♦ fmul_32ns_3i	0	3	143	321
♦ fmul_32ns_3i	0	3	143	321
Memories(0)	0		0	0
> Σ Expressions(67)	0	0	0	848
> Registers(227)			6851	
Channels(0)	0		0	0
> Multiplexers(17)	0		0	1887
> DSP(4)		0		



- Read 2 $M[i][j]$ s and do *fmuls* and *fadds* for them.
- HLS executes two *fadds* in parallel, and so as *fmuls*.

Quantization with AP_Fixed

- Use the “shortest” data type under precision requirements.

```
// mvm.h
#include "ap_fixed.h"
typedef ap_ufixed<8, 5> in_t;
```

Choose an appropriate `out_t`

```
// [xxxxx.xxx] * [xxxxx.xxx] * 64 = [16b.6b]
typedef ap_ufixed<22, 16> out_t;
```

Please try to optimize the design to satisfy the constraints: #DSP < 10, #BRAM < 100.

What's your best latency?

≤ 2000 ?

Optimized design

```
void mvm(in_t M[N][N], in_t V_In[N], out_t V_Out[N]) {
#pragma HLS interface m_axi port=M
#pragma HLS interface m_axi port=V_In
#pragma HLS interface m_axi port=V_Out

#pragma HLS array_partition variable=V_In cyclic factor=16
#pragma HLS array_partition variable=M dim=2 cyclic factor=16
```

```
out_loop:
    for (int i = 0; i < N; i++) {
        out_t sum = 0;
dot_product_loop:
        for (int j = 0; j < N; j+=16) {
#pragma HLS pipeline II=1
            for (int jj = 0; jj < 16; jj++)
#pragma HLS unroll
                sum += V_In[j+jj] * M[i][j+jj];
        }
        V_Out[i] = sum;
    }
}
```

Try this optimized design.

Modules & Loops	Issue Type	Slack	Latency(cycles)	Latency(ns)	Iteration Latency	Interval	Trip Count	Pipelined	BRAM	DSP	FF	LUT
▼ • mvm		-	1222	1.222E4	-	1223	-	no	66	8	18479	24331
▼ ⚙ out_loop		-	1216	1.216E4	19	-	64	no	-	-	-	-
⚙ dot_product_loop		-	9	90.000	7	1	4	yes	-	-	-	-

Design Analysis (1)

```
void mvm(in_t M[N][N], in_t V_In[N], out_t V_Out[N]) {  
#pragma HLS interface m_axi port=M  
#pragma HLS interface m_axi port=V_In  
#pragma HLS interface m_axi port=V_Out
```

```
#pragma HLS array_partition variable=V_In cyclic factor=16  
#pragma HLS array_partition variable=M dim=2 cyclic factor=16
```

```
out_loop:  
    for (int i = 0; i < N; i++) {  
        out_t sum = 0;  
dot_product_loop:  
        for (int j = 0; j < N; j+=16) {  
#pragma HLS pipeline II=1  
            for (int jj = 0; jj < 16; jj++)  
#pragma HLS unroll  
                sum += V_In[j+jj] * M[i][j+jj];  
        }  
        V_Out[i] = sum;  
    }  
}
```

Why array_partition with the configuration
“*cyclic factor=16*”?

To provides enough ports (16) for the
unrolled innermost loop. Finishing reads in
one cycle allows $II=1$.

Design Analysis (2)

```
void mvm(in_t M[N][N], in_t V_In[N], out_t V_Out[N]) {  
#pragma HLS interface m_axi port=M  
#pragma HLS interface m_axi port=V_In  
#pragma HLS interface m_axi port=V_Out  
  
#pragma HLS array_partition variable=V_In cyclic factor=16  
#pragma HLS array_partition variable=M dim=2 cyclic factor=16
```

```
out_loop:  
    for (int i = 0; i < N; i++) {  
        out_t sum = 0;  
        dot_product_loop:  
            for (int j = 0; j < N; j+=16) {  
#pragma HLS pipeline II=1  
                for (int jj = 0; jj < 16; jj++)  
#pragma HLS unroll  
                    sum += V_In[j+jj] * M[i][j+jj];  
            }  
        V_Out[i] = sum;  
    }  
}
```

Why split dot_product_loop into two loops (j and jj)? (**Tiling!**)

To pipeline the partially unrolled loop.
(Equivalent to pipeline + partial unroll.)

Tiling looks more explicit.

Design Analysis (3)

```
void mvm(in_t M[N][N], in_t V_In[N], out_t V_Out[N]) {  
#pragma HLS interface m_axi port=M  
#pragma HLS interface m_axi port=V_In  
#pragma HLS interface m_axi port=V_Out
```


```
#pragma HLS array_partition variable=V_In cyclic factor=16  
#pragma HLS array_partition variable=M dim=2 cyclic factor=16
```

```
out_loop:  
    for (int i = 0; i < N; i++) {  
        out_t sum = 0;  
dot_product_loop:  
        for (int j = 0; j < N; j+=16) {  
#pragma HLS pipeline II=1  
            for (int jj = 0; jj < 16; jj++)  
#pragma HLS unroll  
                sum += V_In[j+jj] * M[i][j+jj];  
        }  
        V_Out[i] = sum;  
    }  
}
```


How does the produced hardware implement
16 additions in the innermost loop?

Look at the *Schedule Viewer*.


Design Analysis (4)




Design Analysis (4)



Design Analysis (4)



Design Analysis (4)



Design Analysis (5)

```
void mvm(in_t M[N][N], in_t V_In[N], out_t V_Out[N]) {  
#pragma HLS interface m_axi port=M  
#pragma HLS interface m_axi port=V_In  
#pragma HLS interface m_axi port=V_Out
```

```
#pragma HLS array_partition variable=V_In cyclic factor=16  
#pragma HLS array_partition variable=M dim=2 cyclic factor=16
```

```
out_loop:  
    for (int i = 0; i < N; i++) {  
        out_t sum = 0;  
dot_product_loop:  
        for (int j = 0; j < N; j+=16) {  
#pragma HLS pipeline II=1  
            for (int jj = 0; jj < 16; jj++)  
#pragma HLS unroll  
                sum += V_In[j+jj] * M[i][j+jj];  
        }  
        V_Out[i] = sum;  
    }  
}
```

How does the produced hardware implement 16 additions in the innermost loop?

Look at the *Schedule Viewer*.

Reduction tree

Why?

- better latency.
- adder is cheap.

