

Parallel Programming for FPGAs (part 2) Practice

Lab1: Discrete Fourier Transform (*DFT*)

- The Discrete Fourier Transform (*DFT*) is a numerical variant of the Fourier Transform.
- The *DFT* transforms a sequence of N complex numbers $\{g_k\} := g_0, g_1, \dots, g_{N-1}$ into another sequence of complex numbers $\{G_n\} := G_0, G_1, \dots, G_{N-1}$, which defined by

$$G_n = \sum_{k=0}^{N-1} g_k \cdot e^{-\frac{i2\pi}{N}kn} = \sum_{k=0}^{N-1} g_k \cdot s^{kn}$$

, where $s = \exp(-\frac{i2\pi}{N})$

Discrete Fourier Transform (2)

$$G_n = \sum_{k=0}^{N-1} g_k \cdot s^{kn}$$

$G = Sg$, where

$$S = \begin{bmatrix} 1 & 1 & 1 & \cdots & 1 \\ 1 & s & s^2 & \cdots & s^{(N-1)} \\ 1 & s^2 & s^4 & \cdots & s^{2(N-1)} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & s^{(N-1)} & s^{2(N-1)} & \cdots & s^{(N-1)(N-1)} \end{bmatrix}$$

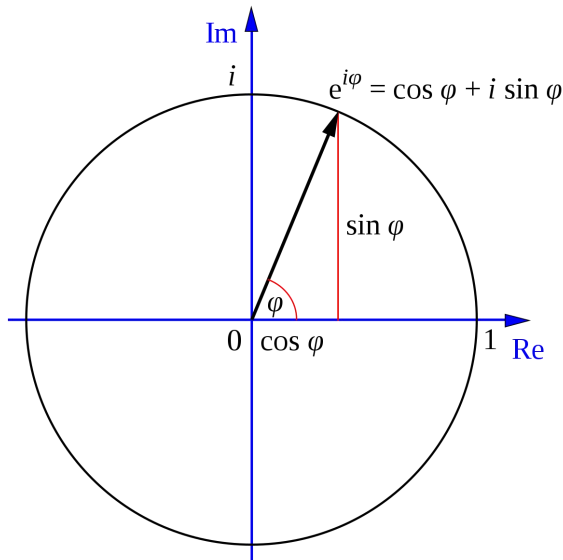
Matrix vector multiplication

How to calculate s^k ($s = \exp(-\frac{i2\pi}{N})$)?

Discrete Fourier Transform (3)

- Euler's Formula: $e^{ix} = \cos x + i \sin x$

$$\begin{aligned} s^k &= \exp\left(-\frac{i2\pi}{N} \cdot k\right) \\ &= \cos\left(-\frac{2\pi}{N} \cdot k\right) + i \sin\left(-\frac{2\pi}{N} \cdot k\right) \end{aligned}$$



$$s^k = s^{k+jN}, j \in N$$

$$s^k = s^{k \bmod N}$$

Summary

Matrix-vector multiplication: $G = Sg$, where

$$S = \begin{bmatrix} 1 & 1 & 1 & \cdots & 1 \\ 1 & s & s^2 & \cdots & s^{(N-1)} \\ 1 & s^2 & s^4 & \cdots & s^{2(N-1)} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & s^{(N-1)} & s^{2(N-1)} & \cdots & s^{(N-1)(N-1)} \end{bmatrix}$$

$$s^k = \exp\left(-\frac{i2\pi}{N} \cdot k\right)$$

$$= \cos\left(-\frac{2\pi}{N} \cdot k\right) + i\sin\left(-\frac{2\pi}{N} \cdot k\right)$$

$$(a + bi) * (c + di) = (ac - bd) + (bc + ad)i$$

Source files

- *dft.h*: typedef, macro, kernel declaration
- *dft.cpp*: HLS kernel definition
- *dft_tb.cpp*: test bench for C-Simulation
- *out.gold.dat*: standard results for input data
- *coefficients1024.h*: $N = 1024$

- $\text{cos_coefficients_table}[k] = \cos\left(-\frac{2\pi}{N} \cdot k\right)$

- $\text{sin_coefficients_table}[k] = \sin\left(-\frac{2\pi}{N} \cdot k\right)$

$$\begin{aligned} s^k &= \exp\left(-\frac{2\pi k i}{N}\right) \\ &= \cos\left(-\frac{2\pi}{N} \cdot k\right) + i \sin\left(-\frac{2\pi}{N} \cdot k\right) \end{aligned}$$

Create baseline

- Create and configure the *HLS* project
- Create a baseline design
 - You may follow the template below.

```
void dft(DTYPE real_in[N], DTYPE imag_in[N],
        DTYPE real_out[N], DTYPE imag_out[N])
{
    // 1. interface pragma, use m_axi protocol
    // 2. define local buffers, and load inputs to them.

    // 3. matrix-vector multiplication loop
    for (int i=0; i<N; i++) {
        // ...
        for (int j=0; j<N; j++) {
            // 3.1. load s[kn] from [cos/sin]_coefficients_table
            // 3.2. calculate complex multiply-accumulate
        }
    }
    // 4. store answers through output interfaces
}
```

run C-Simulation
run C-Synthesis

Create baseline (2)

```
DTYPE real_buf[N], imag_buf[N];
DTYPE real_ans[N], imag_ans[N];
memcpy(real_buf, real_in, N * sizeof(DTYPE));
memcpy(imag_buf, imag_in, N * sizeof(DTYPE));

for (int i=0; i<N; i++) {
    real_ans[i] = imag_ans[i] = 0.;
    int exp = 0;
    for (int j=0; j<N; j++) {
        float real = cos_coefficients_table[exp];
        float imag = sin_coefficients_table[exp];
        real_ans[i] += real * real_buf[j] - imag * imag_buf[j];
        imag_ans[i] += real * imag_buf[j] + imag * real_buf[j];
        exp += i;
        if (exp >= N) exp -= N;
    }
}
memcpy(real_out, real_ans, N * sizeof(DTYPE));
memcpy(imag_out, imag_ans, N * sizeof(DTYPE));
```


Baseline analysis

Modules & Loops	Issue Type	Slack	Latency(cycles)	Latency(ns)	Iteration Latency	Interval	Trip Count	Pipelined	BRAM	DSP	FF	LUT	URAM
▼ • dft	ii II Violation	-	6295599	6.296E7	-	6295600	-	no	20	5	3605	4243	0
Loop 1		-	1025	1.025E4	3	1	1024	yes	-	-	-	-	-
Loop 2		-	1025	1.025E4	3	1	1024	yes	-	-	-	-	-
out_loop_dot_prod_loop	ii II Violation	-	6291469	6.292E7	20	6	1048576	yes	-	-	-	-	-
Loop 4		-	1025	1.025E4	3	1	1024	yes	-	-	-	-	-
Loop 5		-	1025	1.025E4	3	1	1024	yes	-	-	-	-	-

• dft	ii II Violation	-	6295599
Loop 1		-	10.25
Loop 2		-	10.25
out_loop_dot_prod_loop	ii II Violation	-	6291469
Loop 4		-	10.25
Loop 5		-	10.25

- Go To II Violation
- Go To Guidance
- Goto Source

Focus Off

Operation type to include from schedule viewer

Select All

Deselect All

☐ alloca
☐ readreq
☐ br
☐ phi_mux

☐ +
☐ icmp
☐ read
☐ zext

☐ bitcast
☐ getelementptr
☐ write
☐ select

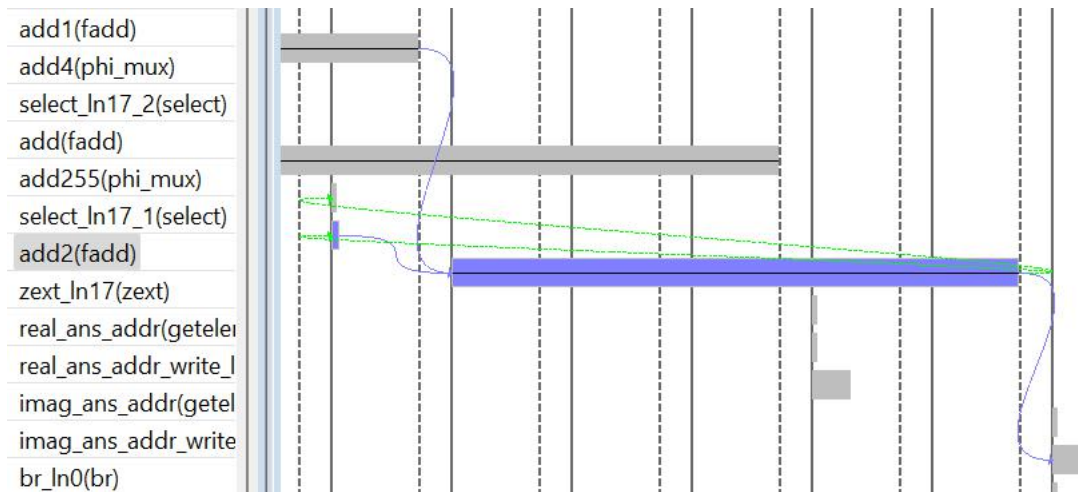
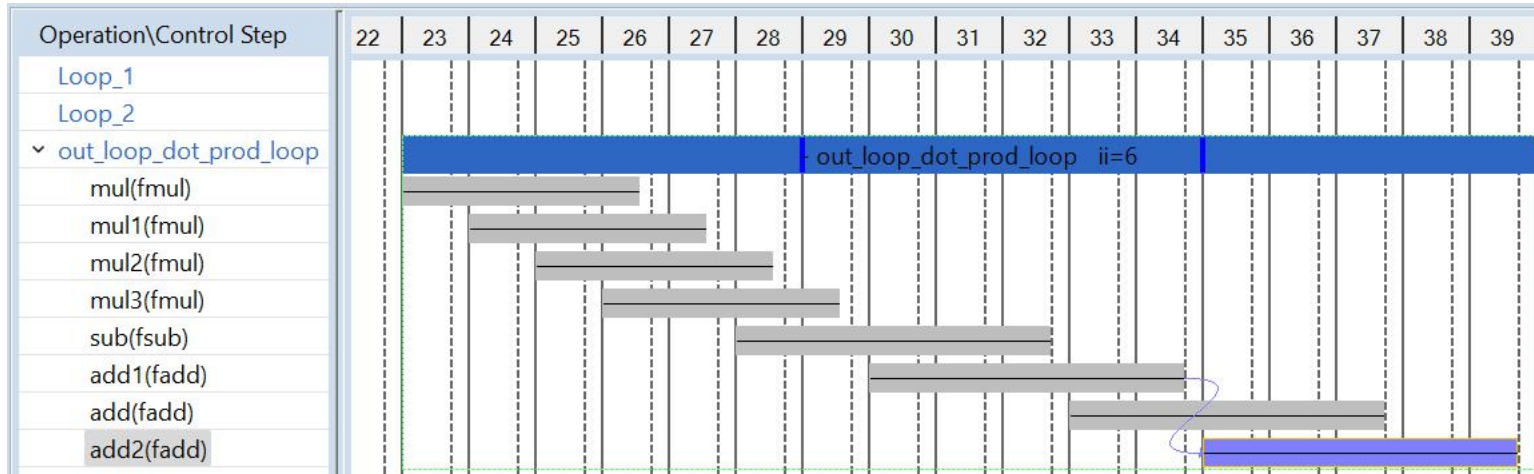
☒ fmul
☐ trunc
☐ partselect
☒ fsub

☒ fadd
☐ writereq
☐ writeresp

Mark cluster(no cluster in this module)

☒ None
☐ Colored
☐ Group

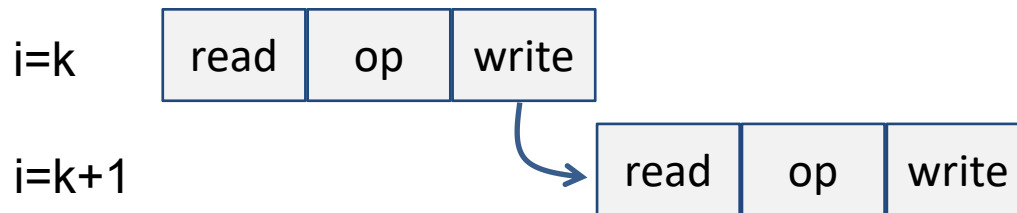
Baseline analysis (2)



Baseline analysis (3)

```
for (int j=0; j<N; j++) {  
    float real = cos_coefficients_table[exp];  
    float imag = sin_coefficients_table[exp];  
    real_ans[i] += real * real_buf[j] - imag * imag_buf[j];  
    imag_ans[i] += real * imag_buf[j] + imag * real_buf[j];  
    exp += i;  
    if (exp >= N) exp -= N;  
}
```

```
for (i) {  
    x = x + a;  
}
```



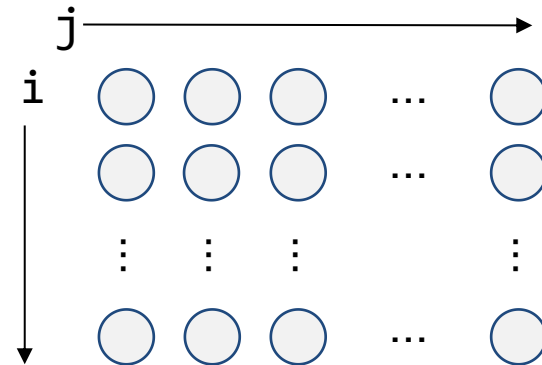
RAW (read-after-write)

$II \geq \text{duration}(\text{read} + \text{op}),$

where $\text{duration}(\text{write})$ can be omitted by forwarding

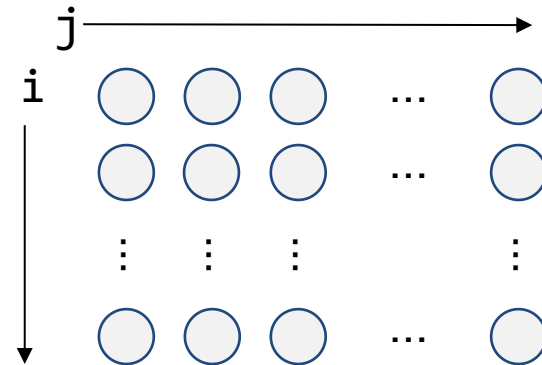
Remove *RAW*

```
for (int i=0; i<N; i++) {  
  ..  
  for (int j=0; j<N; j++) {  
    ..  
    real_ans[i] += ..;  
    imag_ans[i] += ..;  
    ..  
  }  
}
```



```
for (int j=0; i<N; i++) {  
  ..  
  for (int i=0; j<N; j++) {  
    ..  
    real_ans[i] += ..;  
    imag_ans[i] += ..;  
    ..  
  }  
}
```

loop exchange



Apply the optimization. Get design ***dft-opt***.

dft-opt analysis

Modules & Loops	Issue Type	Slack	Latency(cycles)	Latency(ns)	Iteration Latency	Interval	Trip Count	Pipelined	BRAM	DSP	FF	LUT	URAM
▼ • dft		-	1054774	1.055E7	-	1054775	-	no	20	17	4355	5983	0
🔄 Loop 1		-	1025	1.025E4	3	1	1024	yes	-	-	-	-	-
🔄 Loop 2		-	1025	1.025E4	3	1	1024	yes	-	-	-	-	-
🔄 Loop 3		-	1024	1.024E4	1	1	1024	yes	-	-	-	-	-
🔄 Loop 4		-	1024	1.024E4	1	1	1024	yes	-	-	-	-	-
🔄 VITIS_LOOP_33_1_VITIS_LOOP_37_2		-	1048592	1.049E7	18	1	1048576	yes	-	-	-	-	-
🔄 Loop 6		-	1025	1.025E4	3	1	1024	yes	-	-	-	-	-
🔄 Loop 7		-	1025	1.025E4	3	1	1024	yes	-	-	-	-	-

Fully pipelining

Lab2: Fast Fourier Transform(FFT)

- Fast Fourier Transform (*FFT*)
 - DFT implemented by matrix-vector multiplication requires $O(N^2)$ multiplications
 - FFT exploits the structure of the matrix D to reduce the complexity

$$s^{k+\frac{N}{2}} = -s^k$$

$$\begin{aligned} s^k &= \exp\left(-\frac{i2\pi}{N} \cdot k\right) \\ &= \cos\left(-\frac{2\pi}{N} \cdot k\right) + i\sin\left(-\frac{2\pi}{N} \cdot k\right) \end{aligned}$$

- Cooley-Tukey algorithm
 - A divide-and-conquer approach
-

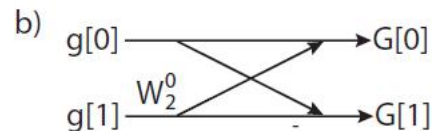
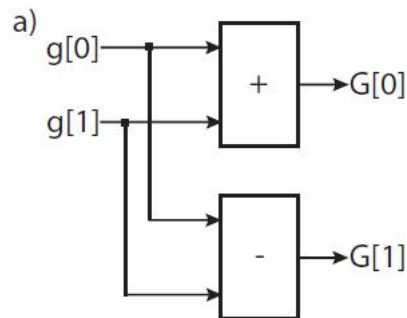
Fast Fourier Transform (2)

- Recall DFT: $X = Dx$
 - D is the DFT coefficients
- For a 2-point DFT

$$S = \begin{bmatrix} W_2^{00} & W_2^{01} \\ W_2^{10} & W_2^{11} \end{bmatrix}, \text{ where } W = e^{-j2\pi}, \quad W_4^{23} = e^{\frac{-j2\pi \cdot 2 \cdot 3}{4}}$$

$$G[0] = g[0] \cdot e^{\frac{-j2\pi \cdot 0 \cdot 0}{2}} + g[1] \cdot e^{\frac{-j2\pi \cdot 0 \cdot 1}{2}} = g[0] + g[1]$$

$$G[1] = g[0] \cdot e^{\frac{-j2\pi \cdot 1 \cdot 0}{2}} + g[1] \cdot e^{\frac{-j2\pi \cdot 1 \cdot 1}{2}} = g[0] - g[1]$$



Fast Fourier Transform (3)

- For a 4-point DFT

$$G[0] = (g[0] + g[2]) + e^{\frac{-j2\pi \cdot 0}{4}}(g[1] + g[3])$$

$$G[1] = (g[0] - g[2]) + e^{\frac{-j2\pi \cdot 1}{4}}(g[1] - g[3])$$

$$G[2] = (g[0] + g[2]) + e^{\frac{-j2\pi \cdot 2}{4}}(g[1] + g[3])$$

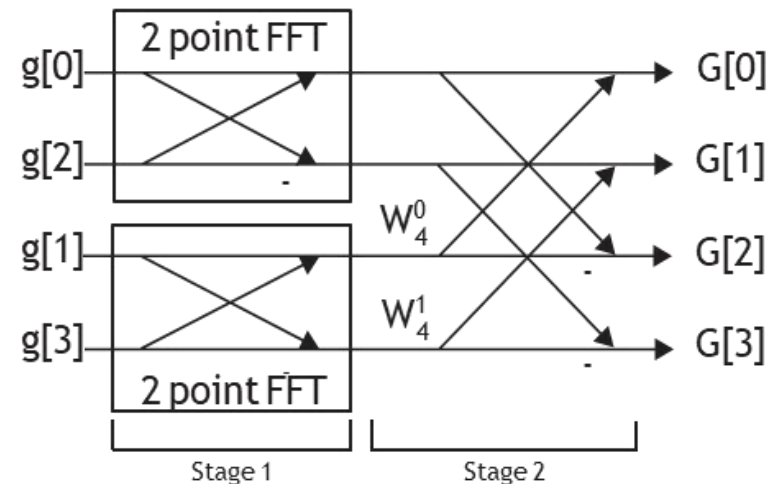
$$G[3] = (g[0] - g[2]) + e^{\frac{-j2\pi \cdot 3}{4}}(g[1] - g[3])$$

$$G[0] = (g[0] + g[2]) + e^{\frac{-j2\pi \cdot 0}{4}}(g[1] + g[3])$$

$$G[1] = (g[0] - g[2]) + e^{\frac{-j2\pi \cdot 1}{4}}(g[1] - g[3])$$

$$G[2] = (g[0] + g[2]) - e^{\frac{-j2\pi \cdot 0}{4}}(g[1] + g[3])$$

$$G[3] = (g[0] - g[2]) - e^{\frac{-j2\pi \cdot 1}{4}}(g[1] - g[3])$$



Fast Fourier Transform (4)

$$G[k] = \sum_{n=0}^{N-1} g[n] \cdot e^{\frac{-j2\pi kn}{N}}, \text{ for } k = 0, \dots, N-1$$

$$\begin{aligned} G[k] &= \sum_{n=0}^{\frac{N}{2}-1} g[2n] \cdot e^{\frac{-j2\pi k(2n)}{N}} + \sum_{n=0}^{\frac{N}{2}-1} g[2n+1] \cdot e^{\frac{-j2\pi k(2n+1)}{N}} \\ &= \sum_{n=0}^{\frac{N}{2}-1} g[2n] \cdot e^{\frac{-j2\pi kn}{N/2}} + e^{\frac{-j2\pi k}{N}} \sum_{n=0}^{\frac{N}{2}-1} g[2n+1] \cdot e^{\frac{-j2\pi kn}{N/2}} \end{aligned}$$

a) $G[k] = A_k + W_N^k B_k$

where A is the DFT result of $\{g_{2n}\} := g_0, g_2, \dots, g_{N-2}$,
and B is the DFT result of $\{g_{2n+1}\} := g_1, g_3, \dots, g_{N-1}$

Fast Fourier Transform (5)

What is $G[k + N/2]$ for $k = 0, \dots, N/2 - 1$?

$$\begin{aligned} G[k + N/2] &= \sum_{n=0}^{\frac{N}{2}-1} g[2n] \cdot e^{\frac{-j2\pi(k+\frac{N}{2})(2n)}{N}} + \sum_{n=0}^{\frac{N}{2}-1} g[2n+1] \cdot e^{\frac{-j2\pi(k+\frac{N}{2})(2n+1)}{N}} \\ &= \sum_{n=0}^{\frac{N}{2}-1} g[2n] \cdot e^{\frac{-j2\pi k(2n)}{N}} \cdot e^{-j2n\pi} + \sum_{n=0}^{\frac{N}{2}-1} g[2n+1] \cdot e^{\frac{-j2\pi k(2n+1)}{N}} \cdot e^{-j(2n+1)\pi} \\ &= \sum_{n=0}^{\frac{N}{2}-1} g[2n] \cdot e^{\frac{-j2\pi kn}{N/2}} - e^{\frac{-j2\pi k}{N}} \sum_{n=0}^{\frac{N}{2}-1} g[2n+1] \cdot e^{\frac{-j2\pi kn}{N/2}} \end{aligned}$$

b) $G[k + N/2] = A_k - W_N^k B_k$

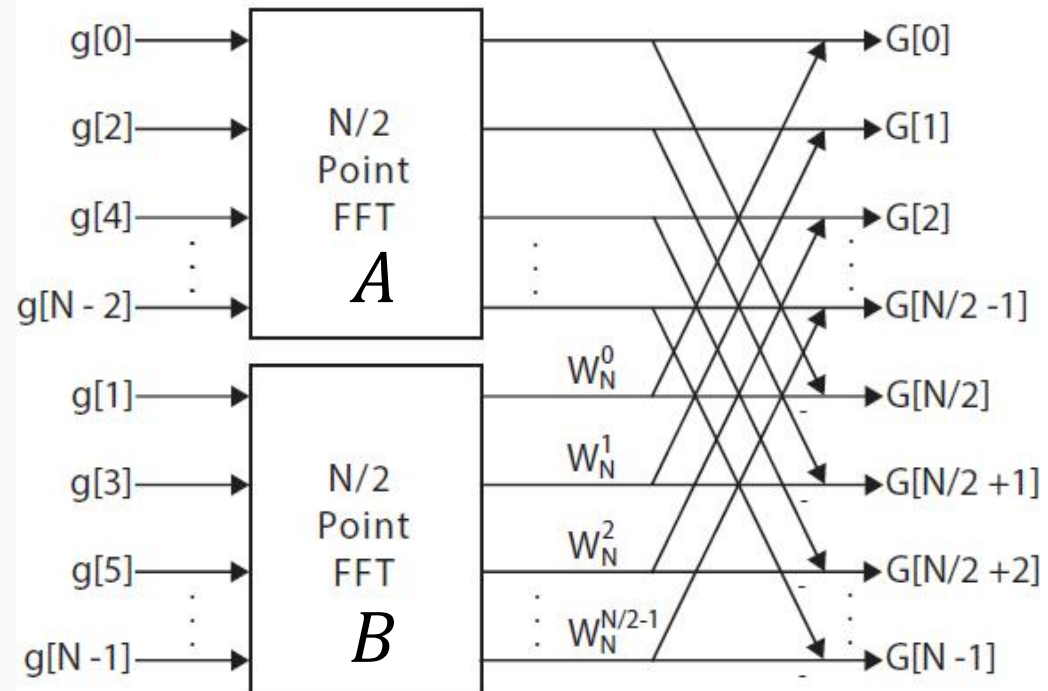
Fast Fourier Transform (6)

a) $G[k] = A_k + W_N^k B_k$

b) $G[k + N/2] = A_k - W_N^k B_k$

```

G0,...,N-1 ← fft_recur(g, N, s):
// DFT of (g0, gs, g2s, ... g(N-1)s)
if N = 1 then G0 ← g0
else
  G0,...,N/2-1 ← fft_recur(g, N/2, 2s)
  GN/2,...,N-1 ← fft_recur(g+s, N/2, 2s)
  for k = 0 to N/2-1 do
    p ← Gk
    q ← exp(-2πik/N) Gk+N/2
    Gk ← p + q
    Gk+N/2 ← p - q
  end for
end if
    
```



Fast Fourier Transform (7)

- Bit reverse: prepare input data layout for *divide-and-conquer*

$(g_0, g_1, \dots, g_{N-1})$ are divided into:
 $(g_0, g_2, \dots, g_{N-2})$ and $(g_1, g_3, \dots, g_{N-1})$.

If the lowest bit of g_n 's original index $(n \& 1)$ is 0, g_k will be divided to the first sequence, otherwise the second sequence.

So, g_k 's new index should have $(n \& 1)$ as its highest bit.

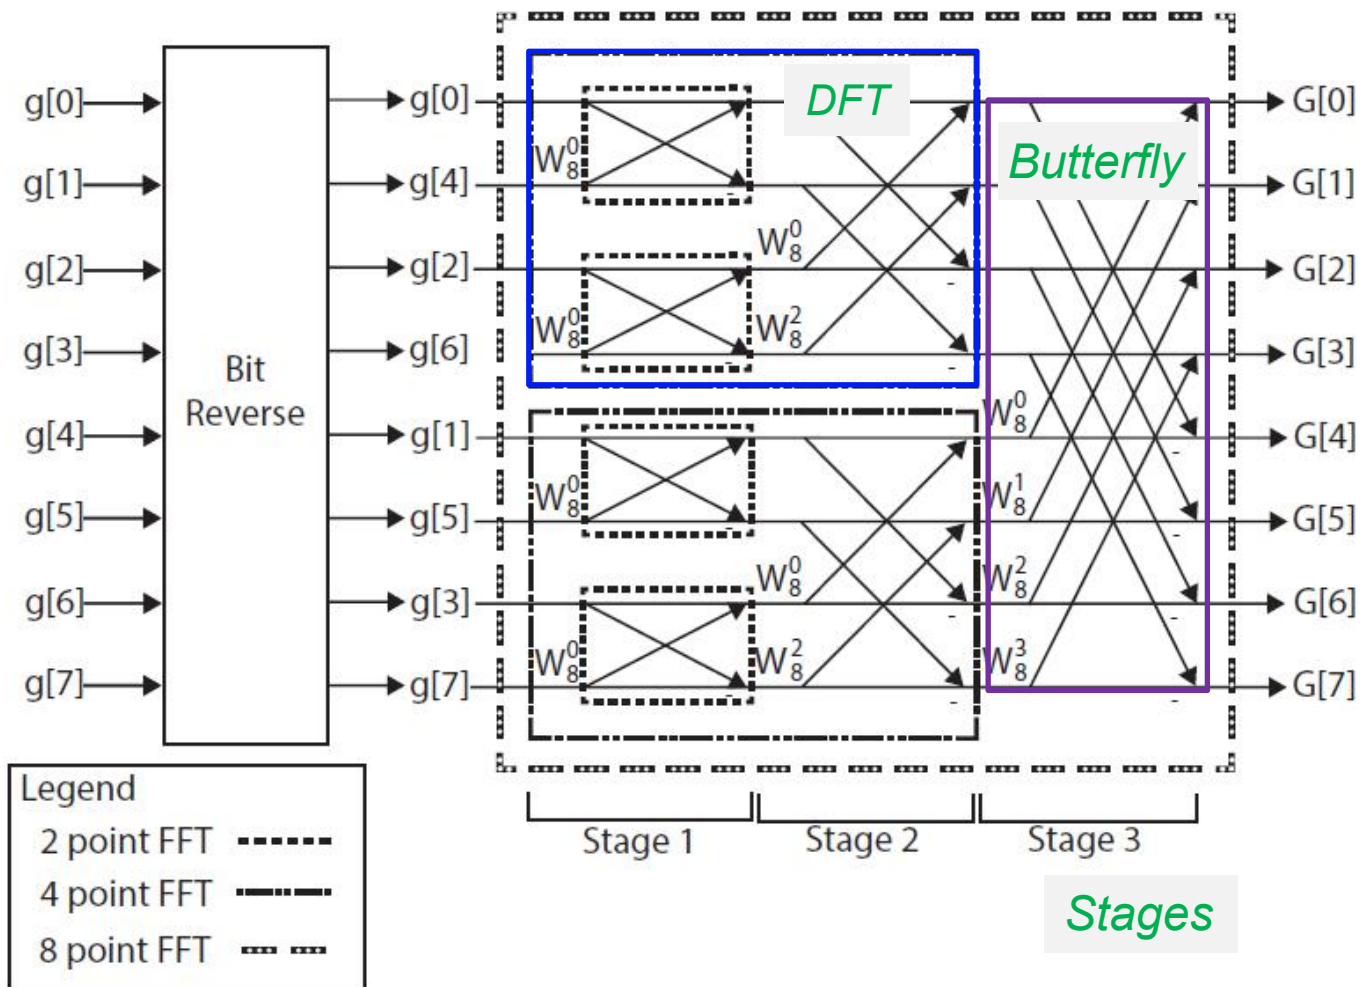
Repeat the procedure, g_n 's new index will be the bit-reverse (*rev*) of its original index.

Index	Binary	Reversed Binary	Reversed Index
0	000	000	0
1	001	100	4
2	010	010	2
3	011	110	6
4	100	001	1
5	101	101	5
6	110	011	3
7	111	111	7

```
rev_index ← rev (index):  
  k = 1  
  rev_index = 0  
  while index > 0 do:  
    rev_index |= (index & 1) << k  
    index = index >> 1  
    k += 1  
  end while
```

Fast Fourier Transform (8)

- For $N = 8$,



Fast Fourier Transform (9)

- Iterative implementation

algorithm bit-reverse-copy(a, A) is
input: Array a of $n=2^{\log_2(n)}$ complex values.
output: Array A of size n .
 $n \leftarrow a.length$
for $k = 0$ **to** $n - 1$ **do**
 $A[\text{rev}(k)] := a[k]$

algorithm iterative-fft is
input: Array a of $n=2^{\log_2(n)}$ complex values.
output: Array A the DFT of a .
bit-reverse-copy(a, A)
 $n \leftarrow a.length$
for $s = 1$ **to** $\log_2(n)$ **do** stage_loop
 $m \leftarrow 2^s$
 $\omega_m \leftarrow \exp(-2\pi i/m)$
 for $k = 0$ **to** $n-1$ **by** m **do** dft_loop
 $\omega \leftarrow 1$
 for $j = 0$ **to** $m/2 - 1$ **do** butterfly_loop
 $t \leftarrow \omega A[k + j + m/2]$
 $u \leftarrow A[k + j]$
 $A[k + j] \leftarrow u + t$
 $A[k + j + m/2] \leftarrow u - t$
 $\omega \leftarrow \omega \cdot \omega_m$
 return A

Summary

- Iterative implementation of the *divide-and-conquer* algorithm

algorithm bit-reverse-copy(a, A) is
input: Array a of $n=2^{\log_2(n)}$ complex values.
output: Array A of size n .
 $n \leftarrow a.\text{length}$
for $k = 0$ **to** $n - 1$ **do**
 $A[\text{rev}(k)] := a[k]$

algorithm iterative-fft is
input: Array a of $n=2^{\log_2(n)}$ complex values.
output: Array A the DFT of a .
bit-reverse-copy(a, A)
 $n \leftarrow a.\text{length}$
for $s = 1$ **to** $\log_2(n)$ **do** stage_loop
 $m \leftarrow 2^s$
 $\omega_m \leftarrow \exp(-2\pi i/m)$
 for $k = 0$ **to** $n-1$ **by** m **do** dft_loop
 $\omega \leftarrow 1$
 for $j = 0$ **to** $m/2 - 1$ **do** butterfly_loop
 $t \leftarrow \omega A[k + j + m/2]$
 $u \leftarrow A[k + j]$
 $A[k + j] \leftarrow u + t$
 $A[k + j + m/2] \leftarrow u - t$
 $\omega \leftarrow \omega \cdot \omega_m$
 return A

Create baseline

- Create and configure the *HLS* project
- Create a baseline design

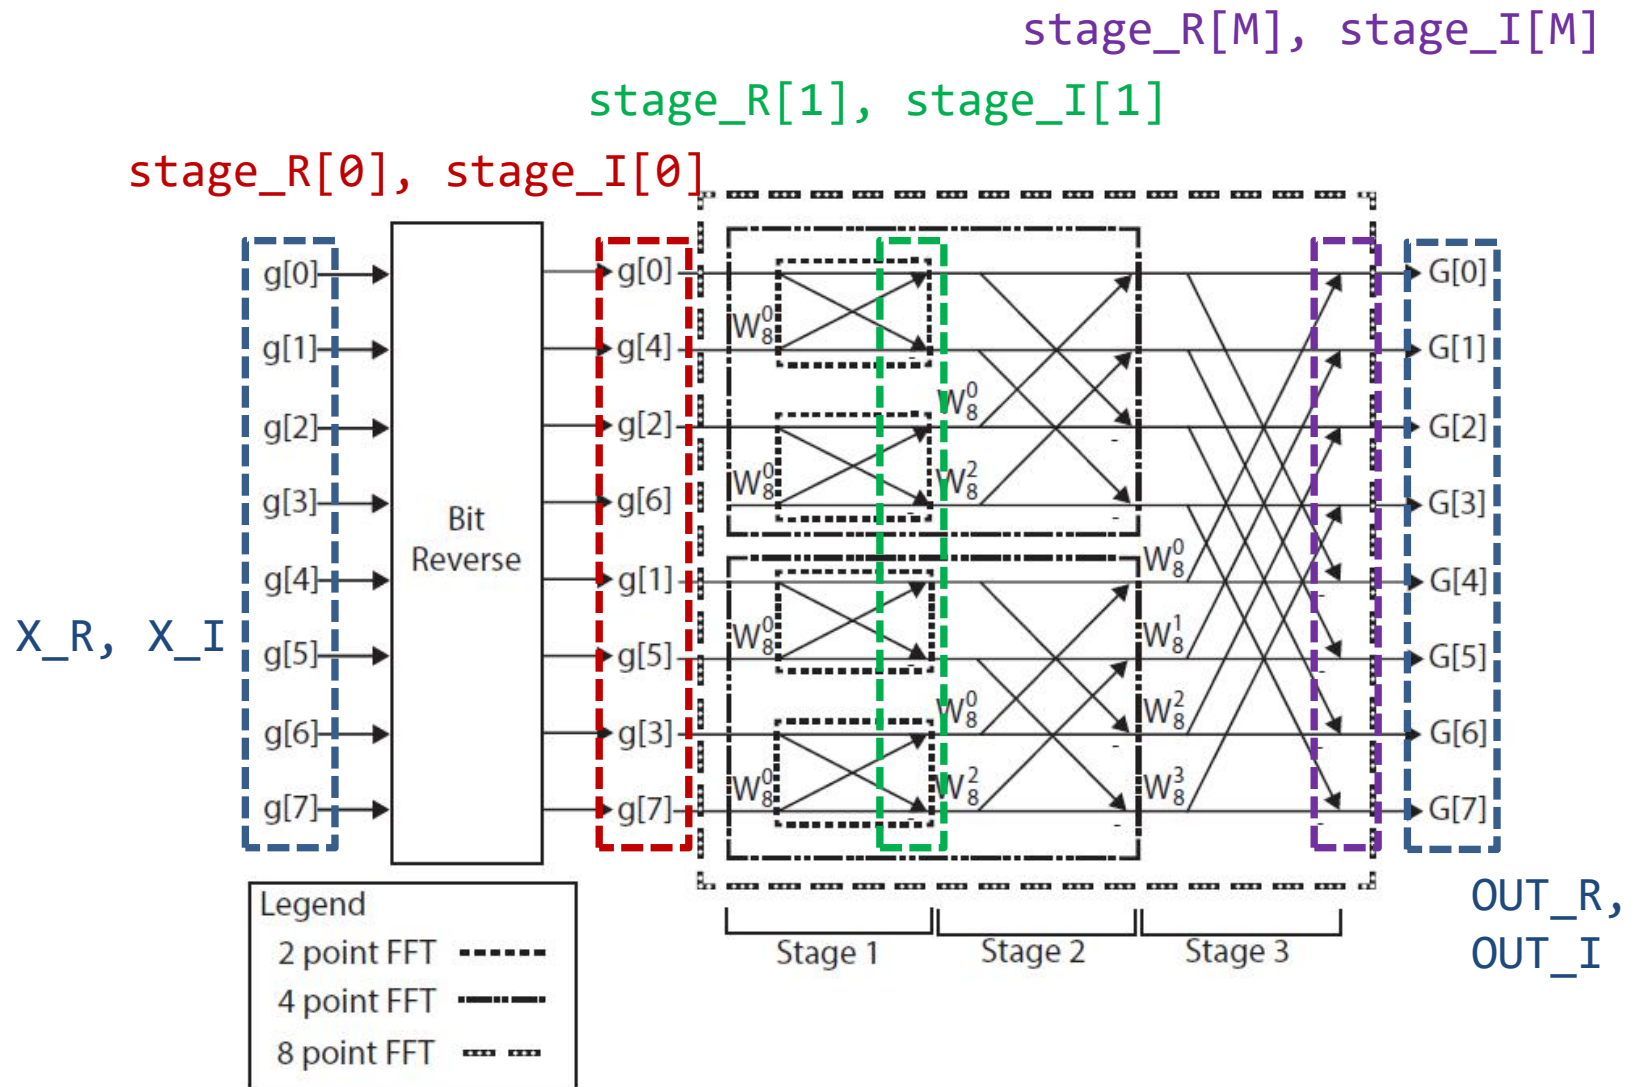
```
void fft(DTYPE X_R[N], DTYPE X_I[N], DTYPE OUT_R[N], DTYPE OUT_I[N])
{
    // interface pragmas, use m_axi protocol

    DTYPE stage_R[M+1][N], stage_I[M+1][N];

    bit_reverse(X_R, X_I, stage_R[0], stage_I[0]);

    for (int stage=0; stage<M; stage++) {
        fft_stage(stage,          stage_R[stage], stage_I[stage],
                        stage_R[stage+1], stage_I[stage+1]);
    }
    store_results(stage_R[M], OUT_R);
    store_results(stage_I[M], OUT_I);
}
```


Create baseline (2)



Create baseline (3)

```
void bit_reverse(DTYPE X_R[N], DTYPE X_I[N],  
                DTYPE OUT_R[N], DTYPE OUT_I[N]) {  
    // Insert your code here  
}
```

algorithm bit-reverse-copy(a, A) is
input: Array a of $N=2^{\log_2(N)}$ complex values.
output: Array A of size N .
 $N \leftarrow a.\text{length}$
 for $i = 0$ **to** $N - 1$ **do**
 $A[\text{rev}(i)] := a[i]$
 end for

how to implement $\text{rev}(k)$?

algorithm $\text{rev}(i)$ is
 $\text{ret} := 0$
 for $k = 0$ **to** $\log_2(N)$ **do**
 $\text{rev} := (\text{rev} \ll 1) \mid ((i \gg k) \& 1)$
 end for
 return ret

```
void bit_reverse(DTYPE X_R[N], DTYPE X_I[N], DTYPE OUT_R[N], DTYPE OUT_I[N]) {  
    for (int i=0; i<N; i++) {  
        int rev = 0;  
        for (int k=0; k<M; k++) {  
            rev = (rev << 1) | ((i >> k) & 1);  
        }  
        OUT_R[rev] = X_R[i];  
        OUT_I[rev] = X_I[i];  
    }  
}
```

Create baseline (4)

```
void fft_stage(int stage, DTYPE X_R[N], DTYPE X_I[N], DTYPE OUT_R[N],
               DTYPE OUT_I[N]) {
    // Insert your code here
}
```

```

 $m \leftarrow 2^s$ 
 $\omega_m \leftarrow \exp(-2\pi i/m)$ 
for  $k = 0$  to  $n-1$  by  $m$  do
     $\omega \leftarrow 1$ 
    for  $j = 0$  to  $m/2 - 1$  do
         $t \leftarrow \omega A[k + j + m/2]$ 
         $u \leftarrow A[k + j]$ 
         $A[k + j] \leftarrow u + t$ 
         $A[k + j + m/2] \leftarrow u - t$ 
         $\omega \leftarrow \omega \cdot \omega_m$ 

```

```

 $\omega_m \leftarrow \exp(-2\pi i/m)$ 
 $= \exp(-2\pi i/N * (N/m))$ 
 $= W\_real[N/m] + i W\_imag[N/m]$ 

```

```

void fft_stage(int stage, DTYPE X_R[N], DTYPE
               X_I[N], DTYPE OUT_R[N], DTYPE OUT_I[N]) {
    int m = 1 << stage, mh = 1 << (stage-1);
    int offset = N >> stage;
    float w_r, w_i, u_r, u_i, a_r, a_i, t_r, t_i;
    for (int k = 0; k < N; k += m) {
        int exp = 0;
        for (int j = 0; j < mh; j++) {
            w_r = W_real[exp], w_i = W_imag[exp];
            u_r = X_R[k+j], u_i = X_I[k+j];
            a_r = X_R[k+j+mh], a_i = X_I[k+j+mh];
            t_r = (w_r * a_r - w_i * a_i);
            t_i = (w_r * a_i + w_i * a_r);
            OUT_R[k+j] = u_r + t_r;
            OUT_I[k+j] = u_i + t_i;
            OUT_R[k+j+mh] = u_r - t_r;
            OUT_I[k+j+mh] = u_i - t_i;
            exp += offset;
        }
    }
}

```

Create baseline (5)

- Run C-Simulation and C-Synthesis

```
void store_results(DTYPE X[N], DTYPE OUT[N]) {  
    memcpy(OUT, X, N * sizeof(DTYPE));  
}
```

Baseline analysis

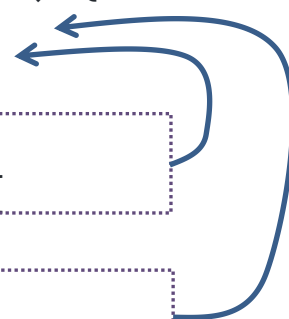
Modules & Loops	Issue Type	Slack	Latency(cycles)	Latency(ns)	Iteration Latency	Interval	Trip Count	Pipelined	BRAM	DSP	FF	LUT	URAM
▼ • fft	ii II Violation	-	-	-	-	-	-	no	76	12	4247	5408	0
🔄 VITIS_LOOP_46_1		-	1025	1.025E4	3	1	1024	yes	-	-	-	-	-
▼ 🔄 VITIS_LOOP_36_1		-	-	-	-	-	10	no	-	-	-	-	-
▼ 🔄 VITIS_LOOP_61_1		-	-	-	1044	-	-	no	-	-	-	-	-
🔄 VITIS_LOOP_63_2	ii II Violation	-	1041	1.041E4	20	2	512	yes	-	-	-	-	-
🔄 Loop 3		-	1025	1.025E4	3	1	1024	yes	-	-	-	-	-
🔄 Loop 4		-	1025	1.025E4	3	1	1024	yes	-	-	-	-	-

undetermined trip count and latency

FFT Optimization

- Inner loops have various behavior

```
stage_loop:
for (s=1; s<=log2(N); s++) {
    m = 1 << s;
    ..
    dft_loop:
    for (k=0; k<N; k+=m) {
        ..
        butterfly_loop:
        for (j=0; j<m/2; j++) {
            ..
        }
    }
}
```



Iterations of `butterfly_loop` and `dft_loop` depend on induction variable “s” of `stage_loop`.

Thus, the `stage_loop` can't be simply *statically* pipelined.

Apply *dataflow* optimization among stages: *dynamic* pipelining

How much stages? $\log_2 N$

Add *dataflow*

```
void fft(DTYPE X_R[N], DTYPE X_I[N], DTYPE OUT_R[N], DTYPE OUT_I[N])
{

#pragma HLS interface m_axi port=X_R
#pragma HLS interface m_axi port=X_I
#pragma HLS interface m_axi port=OUT_R
#pragma HLS interface m_axi port=OUT_I

#pragma HLS dataflow

    DTYPE stage_R[M+1][N], stage_I[M+1][N];
#pragma HLS array_partition variable=stage_R dim=1 complete
#pragma HLS array_partition variable=stage_I dim=1 complete

    bit_reverse(X_R, X_I, stage_R[0], stage_I[0]);


    for (int stage=0; stage<M; stage++) {
        #pragma HLS unroll
        fft_stage(stage+1, stage_R[stage], stage_I[stage], stage_R[stage+1],
stage_I[stage+1]);
    }

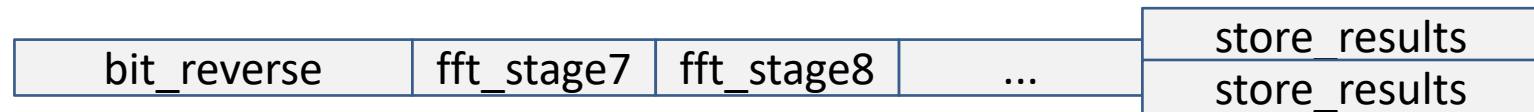
    store_results(stage_R[M], OUT_R);
    store_results(stage_I[M], OUT_I);
}
```

Dataflow analysis

Vitis HLS Console

INFO: [XFORM 203-712] Applying dataflow to function 'fft'

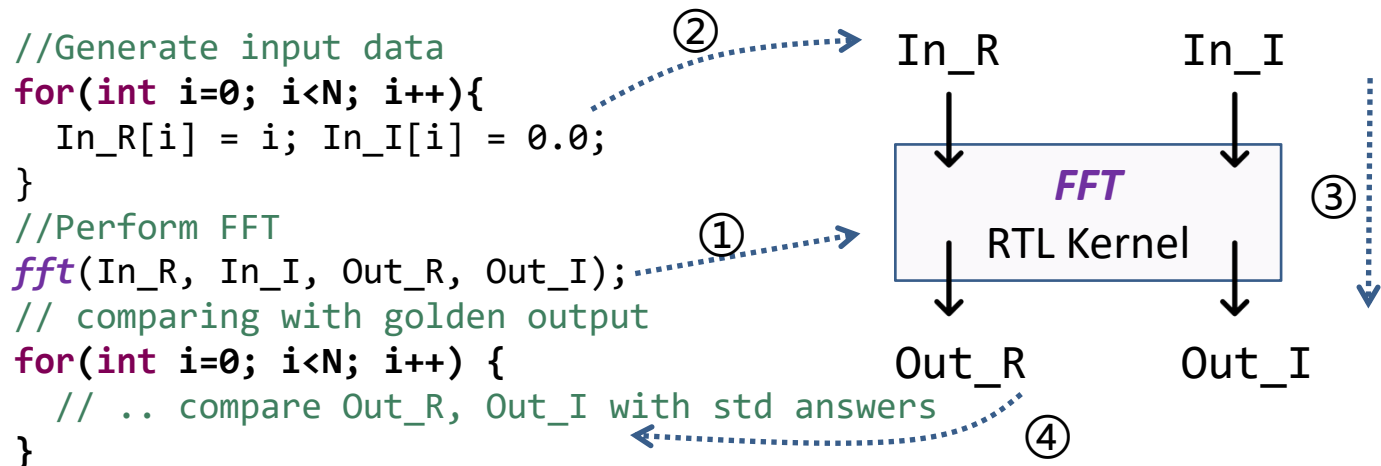
'bit_reverse6'								
'fft_stage7'	Modules & Loops	Issue Type	Slack	Latency(cycles)	Latency(ns)	Iteration Latency	Interval	
'fft_stage8'	▼  fft		-	7374	7.374E4	-	1034	
'fft_stage9'	> ● fft_stage15		-	530	5.300E3	-	530	
'fft_stage10'	> ● fft_stage14		-	530	5.300E3	-	530	
'fft_stage11'	> ● fft_stage13		-	530	5.300E3	-	530	
'fft_stage12'	> ● fft_stage12		-	530	5.300E3	-	530	
'fft_stage13'	> ● fft_stage11		-	530	5.300E3	-	530	
'fft_stage14'	> ● fft_stage10		-	530	5.300E3	-	530	
'fft_stage15'	> ● fft_stage9		-	530	5.300E3	-	530	
'fft_stage16'	> ● fft_stage8		-	530	5.300E3	-	530	
'store_results17'	> ● fft_stage7		-	530	5.300E3	-	530	
'store_results18'.	> ● fft_stage16		-	529	5.290E3	-	529	
	> ● bit_reverse6		-	1033	1.033E4	-	1033	
	> ● store_results17		-	1031	1.031E4	-	1031	
	> ● store_results18		-	1031	1.031E4	-	1031	



Name	Pipelined	Latency	Iteration Latency	Initiation Interval	Trip count
▼ ● fft_stage7	-	530	-	530	-
● VITIS_LOOP_61_1_VITIS_LOOP_63_2	yes	528	18	1	512

C/RTL Cosimulation

- Test register-transfer level (RTL) design synthesized by the HLS tool
 1. Synthesize HLS kernel into RTL
 2. Converts inputs in testbench into input signals
 3. Feed input signals to RTL-level simulation, capture output signals
 4. Continue testbench program with the kernel function replaced by the captured signals.

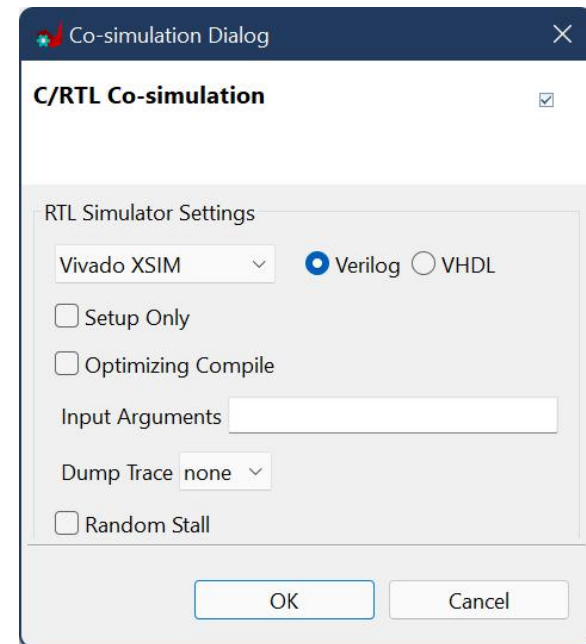
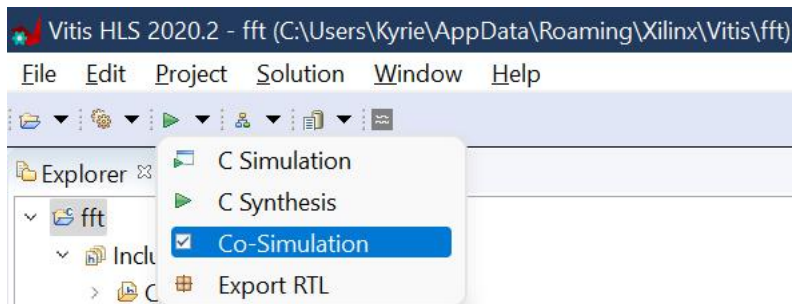


Why Co-Simulation?

- Advantages
 - Reflect any optimizations
 - including architectural opts, such as pipelining
 - Accurate cycle information
 - especially ***stalling*** cycles in *dataflow*
 - Disadvantages
 - cycle-by-cycle simulation is **much slower**
 - limited error messages
-

Co-Simulation

- Run Co-Simulation
 - this is SLOW!



Co-Simulation (2)

Cosimulation Report for 'fft'

General Information

Date: Sat Mar 19 13:02:25 CST 2022

Version: 2020.2 (Build 3064766 on Wed Nov 18 09:12:45 MST 2020)

Project: fft

Status: Pass

Solution: solution1 (Vivado IP Flow Target)

Product family: zynq

Target device: xc7z010-clg400-1

Cosim Options

Tool: Vivado XSIM

RTL: Verilog

Performance Estimates



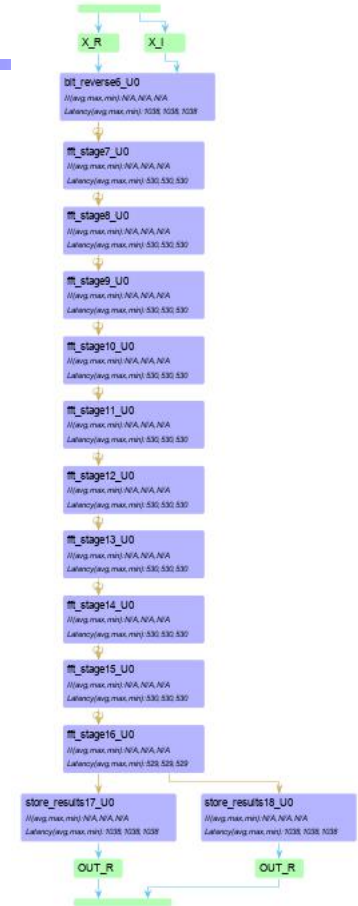
Modules	Avg II	Max II	Min II	Avg Latency	Max Latency	Min Latency
fft				7386	7386	7386
fft_stage15				530	530	530
fft_stage14				530	530	530
fft_stage13				530	530	530
fft_stage12				530	530	530
fft_stage11				530	530	530
fft_stage10				530	530	530
fft_stage9				530	530	530
fft_stage8				530	530	530
fft_stage7				530	530	530
fft_stage16				529	529	529
bit_reverse6				1038	1038	1038
store_results17				1038	1038	1038
store_results18				1038	1038	1038

Co-Simulation (3)

Performance Estimates ?



Modules	Avg II	Max II	Min II	Avg Latency	Max Latency	Min Latency
Open Dataflow Viewer				7386	7386	
Goto Source				530	530	
Open Old Cosimulation Report				530	530	



Properties Warnings Guidance C Source Dataflow

Name	Cosim Category	Cosim Stalling Time	Cosim Read Block Time	Cosim Write Block Time	Cosim Stall No Start	Cosim Stall No Continue	Cosim AVG II	Cosim Max II	Cosim Min II	Cosim AVG Latency	Cosim Max Latency	Cosim Min Latency
bit_reverse6_U0	none	0.00%	0.00%	0.00%	0.00%	0.00%	N/A	N/A	N/A	1038	1038	1038
fft_stage7_U0	none	0.00%	0.00%	0.00%	0.00%	0.00%	N/A	N/A	N/A	530	530	530
fft_stage8_U0	none	0.00%	0.00%	0.00%	0.00%	0.00%	N/A	N/A	N/A	530	530	530
fft_stage9_U0	none	0.00%	0.00%	0.00%	0.00%	0.00%	N/A	N/A	N/A	530	530	530

Better dataflow

Name	BRAM
Channels(22)	88
pingpong	88
stage_I_0_U	4
stage_I_1_U	4
stage_I_10_U	4
stage_I_2_U	4
stage_I_3_U	4
stage_I_4_U	4
stage_I_5_U	4
stage_I_6_U	4
stage_I_7_U	4
stage_I_8_U	4
stage_I_9_U	4
stage_R_0_U	4
stage_R_1_U	4
stage_R_10_U	4
stage_R_2_U	4
stage_R_3_U	4

use *FIFO*?

```
for (int k = 0; k < N; k += m) {
```

```
..
```

```
for (int j = 0; j < mh; j++) {
```

```
..
```

```
u_r = X_R[k+j], u_i = X_I[k+j];  
a_r = X_R[k+j+mh], a_i = X_I[k+j+mh];
```

```
..
```

```
OUT_R[k+j] = u_r + t_r;  
OUT_I[k+j] = u_i + t_i;  
OUT_R[k+j+mh] = u_r - t_r;  
OUT_I[k+j+mh] = u_i - t_i;
```

```
..
```

```
}  
}
```

Not strictly *In-Order*